

## Project Report

This report highlights the Actor and Critic networks with Deep Deterministic Policy Gradients learning agent<sup>i</sup> along with the hyperparameters and provides some references for future work.<sup>ii</sup>

### Actor Network

The network is a feed-forward network with 3 fully connected ReLu activation layers as follows

1. Input 33 (see state size)      => Output 400
2. Input 400                      => Output 300
3. Input 300                      => Output 4 (action size)

with tanh activation function that maps state -> action values.

### Critic Network

The network is a feed-forward network with 4 fully connected leaky ReLu activation layers as follows

1. Input 33 (see state size)      => Output 256
2. Input 256                      => Output 256
3. Input 256                      => Output 128
4. Input 128                      => Output 1

### Learning Agent

The learning agent is created as a class that interacts and learns from the environment with local/target actor and critic network as initialized variants plus Adam variant as optimizer.

### Hyperparameters

- Replay buffer size                      => 1e5
- Minibatch size                      => 128
- Discount factor Gamma                      => 0.99
- TAU for soft update of target parameters      => 1e-3
- Learning rate (Actor)                      => 1e-4
- Learning rate (Critic)                      => 2e-3
- Weight decay                      => 0.0001

### Training

The training of the agent via ddpG function (see Continuous\_Control.ipynb) is running the following steps in every episode:


- Return actions for given state as per current policy from actor/critic local network

- Save experience in replay memory
- Learn every defined update time steps if enough samples are available in memory with random subset to update value parameters using given batch of experience tuples
  1. Get max predicted Q values (for next states) from actor/critic target network
  2. Compute Q targets for current states
  3. Get expected Q values from actor/critic local network
  4. Compute and minimize MSE loss
  5. Soft update target network

In total it is designed to run over 2,000 episodes (with 1,000 iterations per episode) - but it is considered as solved and hence ends if the agent gets an average score of +30 over 100 consecutive episodes.

### Plot of Rewards

Episode 25	Average Score: 1.095	Score: 0.100
Episode 50	Average Score: 1.512	Score: 3.880
Episode 75	Average Score: 2.045	Score: 2.550
Episode 100	Average Score: 2.369	Score: 2.640
Episode 125	Average Score: 3.481	Score: 6.7200
Episode 150	Average Score: 5.127	Score: 14.030
Episode 175	Average Score: 7.248	Score: 15.950
Episode 200	Average Score: 9.712	Score: 19.400
Episode 225	Average Score: 11.839	Score: 17.150
Episode 250	Average Score: 14.217	Score: 26.980
Episode 275	Average Score: 16.768	Score: 20.730
Episode 300	Average Score: 19.248	Score: 23.210
Episode 325	Average Score: 21.664	Score: 17.470
Episode 350	Average Score: 22.689	Score: 32.690
Episode 375	Average Score: 23.552	Score: 21.950
Episode 400	Average Score: 24.105	Score: 37.290
Episode 425	Average Score: 24.456	Score: 26.800
Episode 450	Average Score: 26.390	Score: 34.600
Episode 475	Average Score: 27.398	Score: 33.160
Episode 500	Average Score: 27.755	Score: 15.590
Episode 525	Average Score: 28.852	Score: 32.890
Episode 550	Average Score: 28.292	Score: 30.460
Episode 575	Average Score: 28.635	Score: 33.020
Episode 600	Average Score: 29.733	Score: 37.970
Episode 609	Average Score: 30.036	Score: 34.640
Environment solved in 509 episodes!		Average Score: 30.036



### Future Work

For training purposes, noise adding was paused since the learning process broke down quite early after only a few dozens of episodes – no matter how the different hyperparameters were tuned. In next stage the batch normalization should be added to the respective network architectures to make the learning more robust against added noise.

Additional enhancements can be achieved via implementing some benchmarking from Yan Duan, Xi Chen, Houthooft R., Schulman J., Abbeel P., Benchmarking Deep Reinforcement Learning for Continuous Control, [arXiv:1509.02971](https://arxiv.org/abs/1509.02971).

---

<sup>i</sup> Implementations of the Deep Deterministic Policy Gradients adapted from <https://github.com/udacity/deep-reinforcement-learning/tree/master/ddpg-pendulum> and <https://github.com/udacity/deep-reinforcement-learning/tree/master/ddpg-bipedal>

<sup>ii</sup> Lillicrap et. al., Continuous control with deep reinforcement learning (2015), <https://arxiv.org/abs/1509.02971>